

# PE\_projekt

February 23, 2021

```
[1]: import os
import numpy as np
import torch

import torchvision

from torchvision import datasets, models, transforms
import matplotlib.pyplot as plt

%matplotlib inline
```

```
[2]: # check if CUDA is available
train_on_gpu = False # torch.cuda.is_available()

if not train_on_gpu:
    print('CUDA is not available. Training on CPU ...')
else:
    print('CUDA is available! Training on GPU ...')
```

CUDA is not available. Training on CPU ...

```
[3]: # define training and test data directories
data_dir = 'dataset/'
train_dir = os.path.join(data_dir, 'train/')
test_dir = os.path.join(data_dir, 'test/')
valid_dir = os.path.join(data_dir, 'valid/')

# classes are folders in each directory with these names
classes = ['melanoma', 'nevus', 'seborrheic_keratosis']
```

```
[4]: # load and transform data using ImageFolder

# VGG-16 Takes 224x224 images as input, so we resize all of them
data_transform = transforms.Compose([transforms.Resize(250),
                                     transforms.CenterCrop(224),
                                     transforms.ToTensor()])
train_data = datasets.ImageFolder(train_dir, transform=data_transform)
test_data = datasets.ImageFolder(test_dir, transform=data_transform)
```

```

valid_data = datasets.ImageFolder(valid_dir, transform=data_transform)
# print out some data stats
print('Num training images: ', len(train_data))
print('Num test images: ', len(test_data))
print('Num validation images: ', len(valid_data))

```

```

Num training images: 2000
Num test images: 600
Num validation images: 150

```

```

[5]: # define dataloader parameters
batch_size = 20
num_workers=0

# prepare data loaders
train_loader = torch.utils.data.DataLoader(train_data, batch_size=batch_size,
                                           num_workers=num_workers,
                                           shuffle=True)
test_loader = torch.utils.data.DataLoader(test_data, batch_size=batch_size,
                                           num_workers=num_workers, shuffle=True)
valid_loader = torch.utils.data.DataLoader(valid_data, batch_size=batch_size,
                                           num_workers=num_workers, shuffle=True)

```

```

[6]: # Visualize some sample data

# obtain one batch of training images
dataiter = iter(train_loader)
images, labels = dataiter.next()
images = images.numpy() # convert images to numpy for display

# plot the images in the batch, along with the corresponding labels
fig = plt.figure(figsize=(25, 4))
for idx in np.arange(20):
    ax = fig.add_subplot(2, 20/2, idx+1, xticks=[], yticks=[])
    plt.imshow(np.transpose(images[idx], (1, 2, 0)))
    ax.set_title(classes[labels[idx]])

```

```


/home/lozinske/miniconda3/envs/PE_projekt/lib/python3.7/site-
packages/ipykernel_launcher.py:11: MatplotlibDeprecationWarning: Passing non-
integers as three-element position specification is deprecated since 3.3 and
will be removed two minor releases later.

```

```

# This is added back by InteractiveShellApp.init_path()

```



```
[7]: # Load the pretrained model from pytorch
```

```
vgg16 = models.vgg16(pretrained=True)
```

```
# print out the model structure
```

```
print(vgg16)
```

```
VGG(
```

```
  (features): Sequential(
```

```
    (0): Conv2d(3, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
```

```
    (1): ReLU(inplace=True)
```

```
    (2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
```

```
    (3): ReLU(inplace=True)
```

```
    (4): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1,
```

```
    ceil_mode=False)
```

```
    (5): Conv2d(64, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
```

```
    (6): ReLU(inplace=True)
```

```
    (7): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
```

```
    (8): ReLU(inplace=True)
```

```
    (9): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1,
```

```
    ceil_mode=False)
```

```
    (10): Conv2d(128, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
```

```
    (11): ReLU(inplace=True)
```

```
    (12): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
```

```
    (13): ReLU(inplace=True)
```

```
    (14): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
```

```
    (15): ReLU(inplace=True)
```

```
    (16): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1,
```

```
    ceil_mode=False)
```

```
    (17): Conv2d(256, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
```

```
    (18): ReLU(inplace=True)
```

```
    (19): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
```

```
    (20): ReLU(inplace=True)
```

```
    (21): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
```

```
    (22): ReLU(inplace=True)
```

```
    (23): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1,
```

```
    ceil_mode=False)
```

```
    (24): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
```

```
    (25): ReLU(inplace=True)
```

```
    (26): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
```

```

        (27): ReLU(inplace=True)
        (28): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
        (29): ReLU(inplace=True)
        (30): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1,
ceil_mode=False)
    )
    (avgpool): AdaptiveAvgPool2d(output_size=(7, 7))
    (classifier): Sequential(
      (0): Linear(in_features=25088, out_features=4096, bias=True)
      (1): ReLU(inplace=True)
      (2): Dropout(p=0.5, inplace=False)
      (3): Linear(in_features=4096, out_features=4096, bias=True)
      (4): ReLU(inplace=True)
      (5): Dropout(p=0.5, inplace=False)
      (6): Linear(in_features=4096, out_features=1000, bias=True)
    )
  )
)

```

```
[8]: print(vgg16.classifier[6].in_features)
      print(vgg16.classifier[6].out_features)
```

4096

1000

```
[9]: # Freeze training for all "features" layers
      for param in vgg16.features.parameters():
          param.requires_grad = False
```

```
[10]: import torch.nn as nn

      n_inputs = vgg16.classifier[6].in_features

      # add last linear layer (n_inputs -> 3 classes)
      # new layers automatically have requires_grad = True
      last_layer = nn.Linear(n_inputs, len(classes))

      vgg16.classifier[6] = last_layer

      # if GPU is available, move the model to GPU
      if train_on_gpu:
          vgg16.cuda()

      # check to see that your last layer produces the expected number of outputs
      print(vgg16.classifier[6].out_features)
```

```

[11]: import torch.optim as optim
      # CrossEntropyLoss function is good with dataset with multiple unbalanced
      ↳ classes (different number of samples)
      # specify loss function (categorical cross-entropy)
      criterion = nn.CrossEntropyLoss()

      # specify optimizer (stochastic gradient descent) and learning rate = 0.001,
      ↳ optimize parameters in loss function
      optimizer = optim.SGD(vgg16.classifier.parameters(), lr=0.001) # lr depends on
      ↳ learning results(may differ)

[12]: # number of epochs to train the model
      n_epochs = 20
      valid_loss_min = np.Inf

      for epoch in range(1, n_epochs+1):

          # keep track of training and validation loss
          train_loss = 0.0
          valid_loss = 0.0

          #####
          # train the model #
          #####
          # model by default is set to train
          vgg16.train()
          for batch_i, (data, target) in enumerate(train_loader):
              # move tensors to GPU if CUDA is available
              if train_on_gpu:
                  data, target = data.cuda(), target.cuda()
              # clear the gradients of all optimized variables, to prevent error's
              ↳ accumulation
              optimizer.zero_grad()
              # forward pass: compute predicted outputs by passing inputs to the model
              output = vgg16(data)
              # calculate the batch loss
              loss = criterion(output, target)
              # backward pass: compute gradient of the loss with respect to model
              ↳ parameters
              loss.backward()
              # perform a single optimization step (parameter update)
              optimizer.step()
              # update training loss
              train_loss += loss.item()

              if batch_i % 20 == 19:      # print training loss every specified number
              ↳ of mini-batches

```

```

        print('Training: Epoch %d, Batch %d loss: %.16f' %
              (epoch, batch_i + 1, train_loss / 20))

vgg16.eval() # prep model for evaluation
for batch_i, (data, target) in enumerate(valid_loader):
    # forward pass: compute predicted outputs by passing inputs to the model
    output = vgg16(data)
    # calculate the loss
    loss = criterion(output, target)
    # update running validation loss
    valid_loss += loss.item()

    if batch_i % 20 == 19:      # print training loss every specified number_
    ↪ of mini-batches
        print('Validation: Epoch %d, Batch %d loss: %.16f' %
              (epoch, batch_i + 1, valid_loss / 20))

# print training/validation statistics
# calculate average loss over an epoch
train_loss = train_loss/len(train_loader.dataset)
valid_loss = valid_loss/len(valid_loader.dataset)

print('Epoch: {} \tTraining Loss: {:.6f} \tValidation Loss: {:.6f}'.format(
    epoch+1,
    train_loss,
    valid_loss
))

# save model if validation loss has decreased
if valid_loss <= valid_loss_min:
    print('Validation loss decreased ({:.6f} --> {:.6f}). Saving model ...
    ↪ '.format(
        valid_loss_min,
        valid_loss))
    torch.save(vgg16.state_dict(), 'model.pt')
    valid_loss_min = valid_loss

```

```

Training: Epoch 1, Batch 20 loss: 0.8937232404947281
Training: Epoch 1, Batch 40 loss: 1.7520716816186905
Training: Epoch 1, Batch 60 loss: 2.5959528654813768
Training: Epoch 1, Batch 80 loss: 3.3879931241273882
Training: Epoch 1, Batch 100 loss: 4.1711155086755749
Epoch: 2      Training Loss: 0.041711      Validation Loss: 0.054611
Validation loss decreased (inf --> 0.054611). Saving model ...
Training: Epoch 2, Batch 20 loss: 0.7967709451913834
Training: Epoch 2, Batch 40 loss: 1.5836860701441764
Training: Epoch 2, Batch 60 loss: 2.3688531503081323

```

Training: Epoch 2, Batch 80 loss: 3.1566525831818582  
 Training: Epoch 2, Batch 100 loss: 3.9125271186232569  
 Epoch: 3            Training Loss: 0.039125            Validation Loss: 0.053532  
 Validation loss decreased (0.054611 --> 0.053532). Saving model ...  
 Training: Epoch 3, Batch 20 loss: 0.7570281594991684  
 Training: Epoch 3, Batch 40 loss: 1.4544058024883271  
 Training: Epoch 3, Batch 60 loss: 2.2318968027830124  
 Training: Epoch 3, Batch 80 loss: 2.9838651895523070  
 Training: Epoch 3, Batch 100 loss: 3.7736583530902861  
 Epoch: 4            Training Loss: 0.037737            Validation Loss: 0.050985  
 Validation loss decreased (0.053532 --> 0.050985). Saving model ...  
 Training: Epoch 4, Batch 20 loss: 0.7237058714032173  
 Training: Epoch 4, Batch 40 loss: 1.4230968475341796  
 Training: Epoch 4, Batch 60 loss: 2.1800740018486975  
 Training: Epoch 4, Batch 80 loss: 2.9340826943516731  
 Training: Epoch 4, Batch 100 loss: 3.6843535974621773  
 Epoch: 5            Training Loss: 0.036844            Validation Loss: 0.049986  
 Validation loss decreased (0.050985 --> 0.049986). Saving model ...  
 Training: Epoch 5, Batch 20 loss: 0.7418365031480789  
 Training: Epoch 5, Batch 40 loss: 1.4786423802375794  
 Training: Epoch 5, Batch 60 loss: 2.2104373246431352  
 Training: Epoch 5, Batch 80 loss: 2.9078957885503769  
 Training: Epoch 5, Batch 100 loss: 3.6336459845304487  
 Epoch: 6            Training Loss: 0.036336            Validation Loss: 0.051263  
 Training: Epoch 6, Batch 20 loss: 0.6575327008962631  
 Training: Epoch 6, Batch 40 loss: 1.3752990230917930  
 Training: Epoch 6, Batch 60 loss: 2.0779342621564867  
 Training: Epoch 6, Batch 80 loss: 2.7813130050897596  
 Training: Epoch 6, Batch 100 loss: 3.5476864770054819  
 Epoch: 7            Training Loss: 0.035477            Validation Loss: 0.049243  
 Validation loss decreased (0.049986 --> 0.049243). Saving model ...  
 Training: Epoch 7, Batch 20 loss: 0.6547839269042015  
 Training: Epoch 7, Batch 40 loss: 1.3855033650994302  
 Training: Epoch 7, Batch 60 loss: 2.0509373128414152  
 Training: Epoch 7, Batch 80 loss: 2.7901690840721129  
 Training: Epoch 7, Batch 100 loss: 3.5324532687664032  
 Epoch: 8            Training Loss: 0.035325            Validation Loss: 0.049117  
 Validation loss decreased (0.049243 --> 0.049117). Saving model ...  
 Training: Epoch 8, Batch 20 loss: 0.6893760800361634  
 Training: Epoch 8, Batch 40 loss: 1.4061234891414642  
 Training: Epoch 8, Batch 60 loss: 2.0938444033265116  
 Training: Epoch 8, Batch 80 loss: 2.7429019048810006  
 Training: Epoch 8, Batch 100 loss: 3.4175614401698113  
 Epoch: 9            Training Loss: 0.034176            Validation Loss: 0.049955  
 Training: Epoch 9, Batch 20 loss: 0.6904447585344314  
 Training: Epoch 9, Batch 40 loss: 1.4238117396831513  
 Training: Epoch 9, Batch 60 loss: 2.1013706773519516  
 Training: Epoch 9, Batch 80 loss: 2.7333745867013932

Training: Epoch 9, Batch 100 loss: 3.3881410002708434  
 Epoch: 10            Training Loss: 0.033881            Validation Loss: 0.046973  
 Validation loss decreased (0.049117 --> 0.046973). Saving model ...  
 Training: Epoch 10, Batch 20 loss: 0.6508076488971710  
 Training: Epoch 10, Batch 40 loss: 1.3149775698781014  
 Training: Epoch 10, Batch 60 loss: 1.9995783671736718  
 Training: Epoch 10, Batch 80 loss: 2.6770193666219710  
 Training: Epoch 10, Batch 100 loss: 3.3526929676532746  
 Epoch: 11            Training Loss: 0.033527            Validation Loss: 0.047722  
 Training: Epoch 11, Batch 20 loss: 0.6703696221113205  
 Training: Epoch 11, Batch 40 loss: 1.3369824320077897  
 Training: Epoch 11, Batch 60 loss: 2.0257978051900865  
 Training: Epoch 11, Batch 80 loss: 2.7131043881177903  
 Training: Epoch 11, Batch 100 loss: 3.3446144044399260  
 Epoch: 12            Training Loss: 0.033446            Validation Loss: 0.049403  
 Training: Epoch 12, Batch 20 loss: 0.6251097574830056  
 Training: Epoch 12, Batch 40 loss: 1.3113566055893897  
 Training: Epoch 12, Batch 60 loss: 1.9923828795552254  
 Training: Epoch 12, Batch 80 loss: 2.6851920038461685  
 Training: Epoch 12, Batch 100 loss: 3.2976576119661329  
 Epoch: 13            Training Loss: 0.032977            Validation Loss: 0.049921  
 Training: Epoch 13, Batch 20 loss: 0.6624703913927078  
 Training: Epoch 13, Batch 40 loss: 1.3157240837812423  
 Training: Epoch 13, Batch 60 loss: 1.9309425324201583  
 Training: Epoch 13, Batch 80 loss: 2.5767317861318588  
 Training: Epoch 13, Batch 100 loss: 3.2658273309469221  
 Epoch: 14            Training Loss: 0.032658            Validation Loss: 0.047134  
 Training: Epoch 14, Batch 20 loss: 0.6946130171418190  
 Training: Epoch 14, Batch 40 loss: 1.3513974413275718  
 Training: Epoch 14, Batch 60 loss: 1.9639334157109261  
 Training: Epoch 14, Batch 80 loss: 2.6439152702689173  
 Training: Epoch 14, Batch 100 loss: 3.2515433162450789  
 Epoch: 15            Training Loss: 0.032515            Validation Loss: 0.048121  
 Training: Epoch 15, Batch 20 loss: 0.6201885998249054  
 Training: Epoch 15, Batch 40 loss: 1.3013881355524064  
 Training: Epoch 15, Batch 60 loss: 1.9442445278167724  
 Training: Epoch 15, Batch 80 loss: 2.6474666625261305  
 Training: Epoch 15, Batch 100 loss: 3.2632124155759810  
 Epoch: 16            Training Loss: 0.032632            Validation Loss: 0.047128  
 Training: Epoch 16, Batch 20 loss: 0.6384350180625915  
 Training: Epoch 16, Batch 40 loss: 1.2155491665005684  
 Training: Epoch 16, Batch 60 loss: 1.8687177345156669  
 Training: Epoch 16, Batch 80 loss: 2.4969396248459814  
 Training: Epoch 16, Batch 100 loss: 3.1371933057904244  
 Epoch: 17            Training Loss: 0.031372            Validation Loss: 0.047257  
 Training: Epoch 17, Batch 20 loss: 0.6355775818228722  
 Training: Epoch 17, Batch 40 loss: 1.2600370973348618  
 Training: Epoch 17, Batch 60 loss: 1.8376768201589584



```

Training: Epoch 17, Batch 80 loss: 2.5044110119342804
Training: Epoch 17, Batch 100 loss: 3.1364293128252028
Epoch: 18      Training Loss: 0.031364      Validation Loss: 0.047042
Training: Epoch 18, Batch 20 loss: 0.5969465628266335
Training: Epoch 18, Batch 40 loss: 1.2547573953866960
Training: Epoch 18, Batch 60 loss: 1.8532465308904649
Training: Epoch 18, Batch 80 loss: 2.4716737613081934
Training: Epoch 18, Batch 100 loss: 3.0856217473745344
Epoch: 19      Training Loss: 0.030856      Validation Loss: 0.049140
Training: Epoch 19, Batch 20 loss: 0.6411809772253036
Training: Epoch 19, Batch 40 loss: 1.2447782739996911
Training: Epoch 19, Batch 60 loss: 1.8757026448845864
Training: Epoch 19, Batch 80 loss: 2.4807352736592292
Training: Epoch 19, Batch 100 loss: 3.1147315084934233
Epoch: 20      Training Loss: 0.031147      Validation Loss: 0.046418
Validation loss decreased (0.046973 --> 0.046418). Saving model ...
Training: Epoch 20, Batch 20 loss: 0.5723712176084519
Training: Epoch 20, Batch 40 loss: 1.1699392110109330
Training: Epoch 20, Batch 60 loss: 1.7476233392953873
Training: Epoch 20, Batch 80 loss: 2.3714562132954597
Training: Epoch 20, Batch 100 loss: 3.0541798934340476
Epoch: 21      Training Loss: 0.030542      Validation Loss: 0.047077

```

```

[13]: # track test loss
      # over 3 classes
      test_loss = 0.0
      class_correct = list(0. for i in range(5))
      class_total = list(0. for i in range(5))

      vgg16.eval() # eval mode

      # iterate over test data
      for data, target in test_loader:
          # move tensors to GPU if CUDA is available
          if train_on_gpu:
              data, target = data.cuda(), target.cuda()
          # forward pass: compute predicted outputs by passing inputs to the model
          output = vgg16(data)
          # calculate the batch loss
          loss = criterion(output, target)
          # update test loss
          test_loss += loss.item()*data.size(0)
          # convert output probabilities to predicted class
          _, pred = torch.max(output, 1)
          # compare predictions to true label
          correct_tensor = pred.eq(target.data.view_as(pred))

```

```

        correct = np.squeeze(correct_tensor.numpy()) if 0 else np.
→squeeze(correct_tensor.cpu().numpy()) # not train_on_gpu
        # calculate test accuracy for each object class
        for i in range(batch_size):
            label = target.data[i]
            class_correct[label] += correct[i].item()
            class_total[label] += 1

# calculate avg test loss
test_loss = test_loss/len(test_loader.dataset)
print('Test Loss: {:.6f}\n'.format(test_loss))

for i in range(3):
    if class_total[i] > 0:
        print('Test Accuracy of %5s: %2d%% (%2d/%2d)' % (
            classes[i], 100 * class_correct[i] / class_total[i],
            np.sum(class_correct[i]), np.sum(class_total[i])))
    else:
        print('Test Accuracy of %5s: N/A (no training examples)' % (classes[i]))

print('\nTest Accuracy (Overall): %2d%% (%2d/%2d)' % (
    100. * np.sum(class_correct) / np.sum(class_total),
    np.sum(class_correct), np.sum(class_total)))

```

Test Loss: 0.820256

Test Accuracy of melanoma: 25% (30/117)

Test Accuracy of nevus: 80% (317/393)

Test Accuracy of seborrheic\_keratoses: 36% (33/90)

Test Accuracy (Overall): 63% (380/600)

```

[14]: # obtain one batch of test images
dataiter = iter(test_loader)
images, labels = dataiter.next()
images.numpy()


# move model inputs to cuda, if GPU available
if train_on_gpu:
    images = images.cuda()

# get sample outputs
output = vgg16(images)
# convert output probabilities to predicted class
_, preds_tensor = torch.max(output, 1)
preds = np.squeeze(preds_tensor.numpy()) if 0 else np.squeeze(preds_tensor.
→cpu().numpy()) # not train_on_gpu

```

```
# plot the images in the batch, along with predicted and true labels
fig = plt.figure(figsize=(25, 4))
for idx in np.arange(20):
    ax = fig.add_subplot(2, 20/2, idx+1, xticks=[], yticks=[])
    plt.imshow(np.transpose(images[idx], (1, 2, 0)))
    ax.set_title("{} ({}).format(classes[preds[idx]], classes[labels[idx]]),
                color=("green" if preds[idx]==labels[idx].item() else "red"))
```

/home/lozinske/miniconda3/envs/PE\_projekt/lib/python3.7/site-packages/ipykernel\_launcher.py:19: MatplotlibDeprecationWarning: Passing non-integers as three-element position specification is deprecated since 3.3 and will be removed two minor releases later.



```
[15]: torch.save(vgg16.state_dict(), 'model_backup.pt')
```